Arvind Acharya

+91 8217249312 | arvindacharya.com | LinkedIn | arvindacharya2000@gmail.com | Bengaluru, India

Full Stack AI Engineer with 3+ years building production GenAI and agentic AI systems - LLM agents, RAG, LLM evaluation, and data pipelines in regulated pharma.

SKILLS

Languages: Python, SQL, JavaScript, C++
GenAI & Agentic AI: LLM integration, agentic workflows, multi-agent orchestration, tool calling, MCP, RAG, prompt engineering, context management, LLM evaluation (RAGAS, LLM-as-judge), LLM observability (Langfuse)
Frameworks & Libraries: Microsoft Foundry, Microsoft Agent Framework, Azure AI Search, LangGraph, FastAPI, Streamlit, ChromaDB, sentence-transformers, scikit-learn, React, Copilot Studio, Thunk.AI, C3 AI
Data, Cloud & Tools: Databricks, Snowflake, Denodo, Azure, Docker, Git, Power BI, Power Automate, Aha!, Claude Code, Cursor, GitHub Copilot

EXPERIENCE

GSK (GlaxoSmithKline) *Aug 2023 – Present · Bengaluru*

Full Stack AI Engineer *Jan 2026 – Present*

- Architected and deployed **ARIA**, an LLM-powered agent that automates multi-country employee reimbursements through Workday integration, OCR, and policy-based reasoning, reducing approval time from **10 days to 1 day (90% faster)**.
- Scaled **ARIA** to production as the sole engineer, supporting **4 countries** and processing **1,000+ reimbursement requests per month** via GSK's internal AI platform.
- Identified a production defect on a **live, high-criticality demand forecasting application** — traced the root cause across four user flows and shipped a fix now used **300+ times a month** across a 12-person planning team, **eliminating a recurring manual engineering hand-off**.
- Maintained infrastructure for GSK's agentic AI platform (C3 AI), delivering **12 platform upgrades** across cluster, environment, and application layers with **0 unplanned downtime**.
- Owned end-to-end delivery of **15+ full-stack features** across two production applications using GitHub Copilot and C3 Model Context Protocol (MCP), completing a **10-week roadmap in 3 weeks (70% faster)**.
- Containerized the C3 AI development workflow with Docker and CI/CD pipelines, standardizing development across Windows and macOS and **reducing onboarding time by ~75%** for 10+ engineers.

Data Engineer *Aug 2023 – Dec 2025*

- Built an **LLM-powered PII/PHI classification pipeline** in Databricks for ~75-column clinical datasets, auto-generating three role-based dataset variants; cut processing from **~15 minutes to under 10 seconds** per file and raised classification accuracy from **65% to 98%**.
- Built a multi-layer pipeline moving **2M+ vendor records** from Databricks to Denodo for 4 business teams, and cut runtime by **30%** through Delta optimization and partition pruning.
- Engineered a high-throughput ETL pipeline processing **200K+ daily events** into Snowflake, with automated distribution monitoring that **reduced anomaly detection time by 50%**.
- Applied **semantic similarity analysis** (HuggingFace sentence-transformer embeddings) to R&D investment proposals to flag overlapping initiatives, resulting in **4 proposal restructures**; automated the allocation process, **reducing data quality issues to 0**.
- Built reporting automation that cut data consolidation from **2 hours to under 1 minute**, a portfolio health dashboard used by all product owners and directors (**1 day to under 30 seconds**), and a Concur-based spend tracker that **doubled budget forecasting accuracy**.

PROJECTS

Docs RAG Assistant | *RAG · LangGraph · ChromaDB · RAGAS · Langfuse · FastAPI* GitHub

- Built a RAG assistant with a **LangGraph router** (documentation retrieval, live web search, or decline), a ChromaDB vector store, header-aware chunking, and cross-encoder reranking, served through a FastAPI backend deployed on Render.
- Built an **LLM evaluation pipeline** - RAGAS scoring with Claude Haiku as an independent judge over a hand-written golden set - reaching **95.6% faithfulness, 85% retrieval hit-rate, and 100% routing accuracy**; added Langfuse tracing with quality and cost dashboards and user feedback wired into scores.
- Fixed a corpus-skew bug (one file was **45.7% of the vector store**) and 4 deployment memory issues, cutting a **905MB import to 0MB**.

RegBrief | *Agentic AI · Tool calling · Claude API* GitHub

- Built a **multi-step agent** (plan → parallel tool search → draft → self-critique → finalize) using native tool calling across PubMed, ClinicalTrials.gov, and openFDA, producing source-traceable regulatory briefs with Word/PDF export.
- Cut cost per brief by **~40% ($0.50 → $0.30)** with prompt caching and capped tool results; validated across **4 distinct regulatory pathways with zero code changes**; fixed 3 control-flow bugs with mocked-API unit tests.

Combat24 | *Full-Stack Developer* *combat24.in*

- Reframed a redesign request as a lead-generation problem and rebuilt the site SEO-first, growing organic enquiries from **1–2 to ~10 per week (5x)**.
- Built a full-stack membership platform - member app, admin portal, and QR-based access system - **replacing third-party software** and supporting daily operations for **200+ members**.

EDUCATION

BTech, Computer Science and Engineering | Ramaiah University of Applied Sciences | CGPA: 8.7 *Jul 2019 – May 2023*